
Capstone Lab: End-to-End Automated Deployment Pipeline

In this capstone project, you will design and implement a complete end-to-end automated deployment pipeline that demonstrates mastery of modern CI/CD practices. You'll build a scalable, secure, and monitored web application infrastructure using Infrastructure as Code, containerization, and CI/CD pipelines.

This capstone integrates multiple technologies and best practices to create a production-ready system that can automatically deploy.

In this capstone, you will:

- Design and implement Infrastructure as Code using Ansible
- Containerize applications using Docker and orchestrate with Docker Compose
- Build comprehensive CI/CD pipelines with automated testing and deployment
- Implement security best practices and compliance checking
- Demonstrate knowledge of modern deployment strategies

Estimated completion time

130 minutes

Task 1: Project Setup and Infrastructure Design

In this task, you will set up the capstone project structure, design the infrastructure architecture, and create the foundational configuration files.

Step 1: Create the capstone project structure

Create a folder named **Capstone-CICD** on the Desktop. Open **Capstone-CICD** in **Visual Studio Code**.

Step 2: Initialize project structure

Open VSCode terminal (Terminal → New Terminal) and create the complete project structure:

```
# App
```

```
mkdir -p application/backend/app application/backend/tests
```

```
# Ansible
```

```
mkdir -p infrastructure/ansible/{inventories,group_vars,playbooks}
```

```
# Helper script for docker install (copy from Github Docker repo)
```

```
curl -fsSL https://get.docker.com -o get-docker.sh
```

```
cp ~/get-docker.sh infrastructure/ansible/get-docker.sh 2>/dev/null || true
```

```
# Workflows
```

```
mkdir -p .github/workflows
```

Step 3: Python virtual environment & app code

Create the Virtual Environment:

```
cd ~/Desktop/Capstone-CICD/application/backend
```

```
python3 -m venv .venv
```

```
source .venv/bin/activate
```

Create the requirements file:

application/backend/requirements.txt

```
Flask==2.3.3
gunicorn==21.2.0
pytest==7.4.3
pytest-html==3.2.0
pytest-cov==4.1.0
```

Create the application code file:

application/backend/app/app.py

```
from flask import Flask, jsonify

app = Flask(__name__)

@app.get("/")
def home():
    return "Hello from Capstone CI/CD!", 200

@app.get("/api/health")
def health():
    return jsonify({"status": "healthy", "version": "latest"}), 200

if __name__ == "__main__":
    app.run(host="0.0.0.0", port=5000)
```

Course Title Deliverable Type

Testing the app:

Run the following command to install all testing dependencies:

```
pip install pytest pytest-cov pytest-html flask
```

This will install:

pytest → testing framework

pytest-cov → for coverage reports

pytest-html → for generating HTML reports

Create the test code files:

application/backend/tests/conftest.py

- This is common in Flask/Django/WSGI entry points.
- It guarantees that when the app runs (especially in Docker, Gunicorn, or other environments), Python knows where your app package is, even if the current working directory is different.

```
import os, sys
# Ensure we can import the `app` package
BASE_DIR = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
sys.path.insert(0, BASE_DIR)
```

application/backend/tests/test_app.py

```
from app.app import app

def test_home():
    client = app.test_client()
    rv = client.get("/")
    assert rv.status_code == 200
    assert b"Hello" in rv.data
```

```
def test_health():
    client = app.test_client()
    rv = client.get("/api/health")
    assert rv.status_code == 200
```

```
data = rv.get_json()  
assert data["status"] == "healthy"
```

Run tests locally:

```
cd ~/Desktop/Capstone-CICD/application/backend  
source .venv/bin/activate  
pytest -v --maxfail=1 --disable-warnings \  
--html=report.html --self-contained-html \  
--cov=app --cov-report=html
```

You should see 2 tests passing. If you get ModuleNotFoundError: No module named 'app', the conftest.py above fixes it.

```
● (.venv) student@labuser-virtual-machine:~/Desktop/Capstone/application/backend$ pytest -  
v --maxfail=1 --disable-warnings --html=report.html --self-contained-html --cov=app  
--cov-report=html  
===== test session starts =====  
platform linux -- Python 3.10.12, pytest-8.4.2, pluggy-1.6.0 -- /home/student/Desktop/Ca  
pstone/application/backend/.venv/bin/python3  
cachedir: .pytest_cache  
metadata: {'Python': '3.10.12', 'Platform': 'Linux-6.5.0-44-generic-x86_64-with-glibc2.3  
5', 'Packages': {'pytest': '8.4.2', 'pluggy': '1.6.0'}, 'Plugins': {'metadata': '3.1.1',  
'html': '4.1.1', 'cov': '7.0.0'}}  
rootdir: /home/student/Desktop/Capstone/application/backend  
plugins: metadata-3.1.1, html-4.1.1, cov-7.0.0  
collected 2 items  
  
tests/test_app.py::test_home PASSED [ 50%]  
tests/test_app.py::test_health PASSED [100%]  
  
===== tests coverage =====  
coverage: platform linux, python 3.10.12-final-0
```

Task 2: Dockerize the application

Step 1: Create Docker configuration File

application/backend/Dockerfile

```
# Production image
FROM python:3.11-slim

ENV PYTHONDONTWRITEBYTECODE=1 \
    PYTHONUNBUFFERED=1

WORKDIR /app

# Install system deps (curl helps healthcheck; optional)
RUN apt-get update && apt-get install -y --no-install-recommends curl && \
    rm -rf /var/lib/apt/lists/*

# Copy and install
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Copy app code
COPY app/ ./app/

EXPOSE 5000

# Healthcheck (tries /api/health)
HEALTHCHECK --interval=30s --timeout=5s --retries=3 \
CMD curl -fsS http://127.0.0.1:5000/api/health || exit 1

# Gunicorn entrypoint
CMD ["gunicorn", "--bind", "0.0.0.0:5000", "app.app:app", "--workers", "2", "--timeout", "30"]
```

Step 7: Build and run locally

```
cd ~/Desktop/Capstone-CICD/application/backend
```

```
docker build -t capstone-flask:local .
```

```
=> [1/6] FROM docker.io/library/python:3.11-slim@sha256:9bffe4353b925a1656688797 4.3s
=> => resolve docker.io/library/python:3.11-slim@sha256:9bffe4353b925a1656688797 0.0s
=> => sha256:44350d10c02e7ab437e3fe5a05e3405115ece5972b2b9f7cd0d 1.29MB / 1.29MB 3.3s
=> => sha256:4dc2c3222cdbf7b5e9d5c68653d42c7289ddf2bfaa17b12c9 14.64MB / 14.64MB 1.2s
=> => sha256:b25238518c0cca0928b2117b90cee455c3fbdb7d605f92131e5cc92 249B / 249B 1.4s
=> => sha256:9bffe4353b925a1656688797ebc68f9c525e79b1d377a764d 10.37kB / 10.37kB 0.0s
=> => sha256:70f7abeaf1577b30229dd1d7784d6c053a29104a56bb353fe23 1.75kB / 1.75kB 0.0s
=> => sha256:bf02a2b853727373d9065ccd2cc7d40df56d6f1b8256ae5f361 5.38kB / 5.38kB 0.0s
=> => extracting sha256:44350d10c02e7ab437e3fe5a05e3405115ece5972b2b9f7cd0d68d23 0.1s
=> => extracting sha256:4dc2c3222cdbf7b5e9d5c68653d42c7289ddf2bfaa17b12c96101475 0.9s
=> => extracting sha256:b25238518c0cca0928b2117b90cee455c3fbdb7d605f92131e5cc92f 0.0s
=> [internal] load build context 0.0s
=> => transferring context: 1.24kB 0.0s
=> [2/6] WORKDIR /app 0.1s
=> [3/6] RUN apt-get update && apt-get install -y --no-install-recommends curl 27.0s
=> [4/6] COPY requirements.txt .
=> [5/6] RUN pip install --no-cache-dir -r requirements.txt 20.3s
=> [6/6] COPY app/ ./app/ 0.0s
=> exporting to image 0.2s
=> => exporting layers 0.2s
=> => writing image sha256:6e3066a50e9dbe8b430cd7f0a09d85f5ae9306d93b109beb7235c 0.0s
=> => naming to docker.io/library/capstone-flask:local 0.0s
```

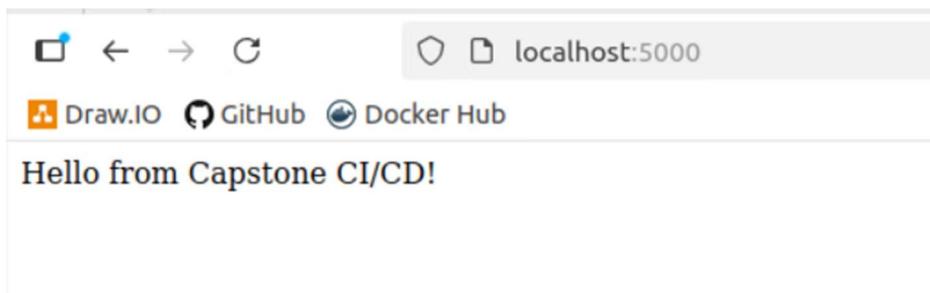
Docker login ??

```
docker run --rm -d -p 5000:5000 --name capstone_flask capstone-flask:local
```

- (.venv) **student@labuser-virtual-machine:~/Desktop/Capstone/application/backend\$** docker run --rm -d -p 5000:5000 --name capstone_flask capstone-flask:local bb689d117fc3cb358cc0e69c0002d0253b52fd386d31df8a1204604b8b1d3a2f
- (.venv) **student@labuser-virtual-machine:~/Desktop/Capstone/application/backend\$** █

```
curl -fsS http://localhost:5000/api/health
```

- (.venv) **student@labuser-virtual-machine:~/Desktop/Capstone/application/backend\$** curl -fsS http://localhost:5000/api/health {"status": "healthy", "version": "latest"}
- (.venv) **student@labuser-virtual-machine:~/Desktop/Capstone/application/backend\$** █



```
docker stop capstone_flask
```

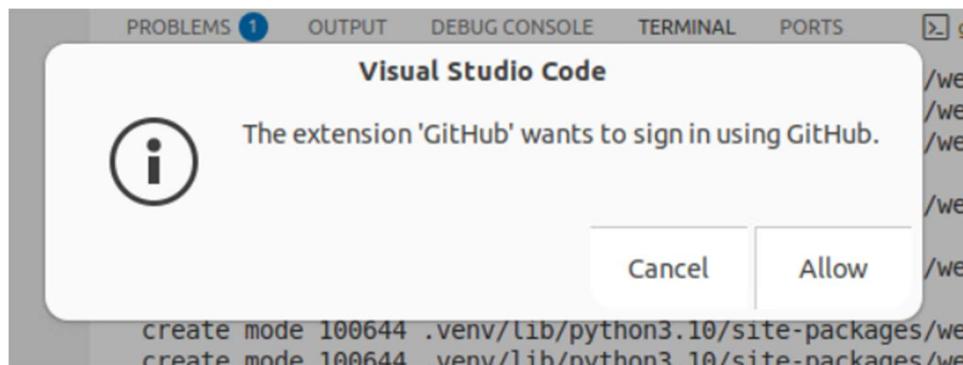
Task 3: Initialize Git and push to GitHub

First, create a new repo on GitHub named **capstone-cicd** (empty i.e. no README).

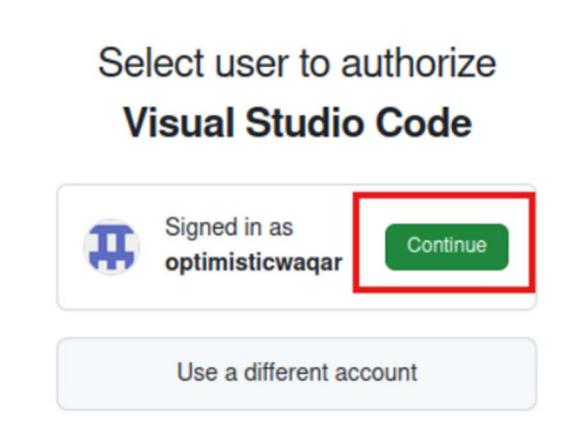
Then:

```
cd ~/Desktop/Capstone-CICD  
git init  
git branch -M main  
git remote add origin https://github.com/<YOUR_USER>/capstone-cicd.git  
git add .  
git commit -m "Initial commit: app, tests, Dockerfile, ansible skeleton, workflows scaffolding"  
git push -u origin main
```

Allow the extension 'Github' to sign in to Github.



Click continue:



Task 4: GHCR (GitHub Container Registry) setup

You have two choices:

- **Public image (easiest):** No auth required to pull
- **Private image: Requires login (GHCR_USER, GHCR_PAT)**

(A) Make image Public

Do nothing now; the workflow tags will create it, and you can change visibility to Public from the package page later.

(B) Keep image Private

Create a **Personal Access Token (classic)** with scopes:

- read:packages, write:packages, (optional delete:packages), repo

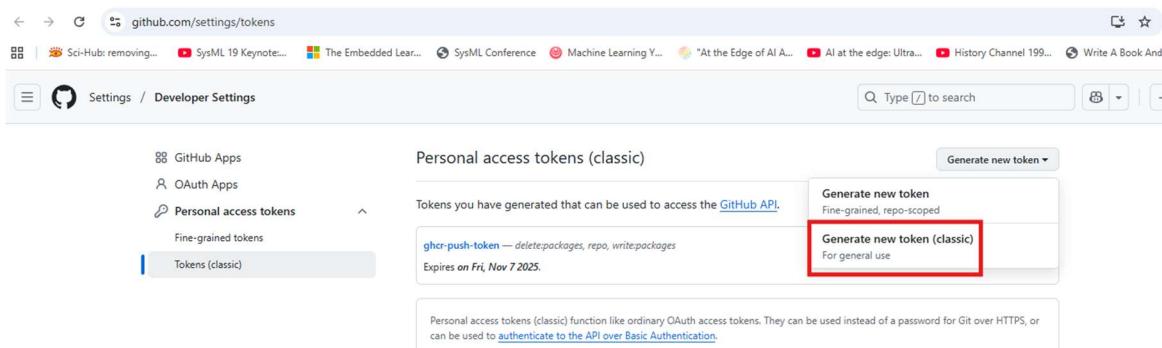
Then add repo **Secrets** (GitHub → Repo → Settings → Secrets and variables → Actions):

- GHCR_USER = your GitHub username (e.g., optimisticwaqar)
- GHCR_PAT = your personal access token value

Steps to Create a GHCR PAT

Step 1: Go to your GitHub Settings

- Open <https://github.com/settings/tokens>.
- Click “**Generate new token**” → “**Generate new token (classic)**”.



Step 2: Give the Token a Name & Expiration

- Add a descriptive name like **GHCR_PAT**.
- Set an expiration date (or “No expiration” if needed for CI/CD).

New personal access token (classic)

Personal access tokens (classic) function like ordinary OAuth access tokens. They can be used instead of a password for Git over HTTPS, or can be used to [authenticate to the API over Basic Authentication](#).

Note

GHCR_PAT

What's this token for?

Expiration

No expiration ▾

⚠ GitHub strongly recommends that you set an expiration date for your token to help keep your information secure. [Learn more](#)

Step 3: Select Required Scopes

For GHCR, you typically need:

- read:packages → if only pulling images
- write:packages → if pushing images
- delete:packages → if you want to delete packages
- repo → if the package is private and linked to a repo

Select scopes

Scopes define the access for personal tokens. [Read more about OAuth scopes.](#)

<input checked="" type="checkbox"/> repo	Full control of private repositories
<input checked="" type="checkbox"/> repo:status	Access commit status
<input checked="" type="checkbox"/> repo_deployment	Access deployment status
<input checked="" type="checkbox"/> public_repo	Access public repositories
<input checked="" type="checkbox"/> repo:invite	Access repository invitations
<input checked="" type="checkbox"/> security_events	Read and write security events
<input type="checkbox"/> workflow	Update GitHub Action workflows
<input checked="" type="checkbox"/> write:packages	Upload packages to GitHub Package Registry
<input checked="" type="checkbox"/> read:packages	Download packages from GitHub Package Registry
<input checked="" type="checkbox"/> delete:packages	Delete packages from GitHub Package Registry

For CI/CD pushing: **check** write:packages **and** read:packages (and repo if your repo is private).

Step 4: Generate the Token

- Click Generate token.
- Copy it immediately (you can't see it later).

The screenshot shows the GitHub Developer Settings page at <https://github.com/settings/tokens>. The user is in the 'Personal access tokens (classic)' section. A message at the top states: "Some of the scopes you've selected are included in other scopes. Only the minimum set of necessary scopes has been saved." Below this, there is a note: "Tokens you have generated that can be used to access the [GitHub API](#)." A warning message in a blue box says: "Make sure to copy your personal access token now. You won't be able to see it again!" A specific token, "ghp_wmLaqhj6DcggHY3o2MvwFq8bmgW5qV4GJwkP", is highlighted with a red box. This token has the scope "delete_packages, repo, write_packages" and was last used "within the last week". It also has a "Delete" button. A note at the bottom explains: "Personal access tokens (classic) function like ordinary OAuth access tokens. They can be used instead of a password for Git over HTTPS, or can be used to [authenticate to the API over Basic Authentication](#)."

Task 5: Ansible: inventory, vars, deploy playbook

Step 1: Create development.yml

infrastructure/ansible/inventories/development.yml

```
---
all:
  hosts:
    localhost:
      ansible_connection: local
      ansible_python_interpreter: /usr/bin/python3
  vars:
    app_port: 5000
    image_ref: "ghcr.io/<YOUR_USER>/capstone-cicd/capstone-flask:latest"
    ghcr_private: false # set true if your GHCR package is private
```

Note: Replace <YOUR_USER> with your GitHub username.

Step 2: Create all.yml

infrastructure/ansible/group_vars/all.yml

```
---
app_port: 5000
image_ref: "ghcr.io/<YOUR_USER>/capstone-cicd/capstone-flask:latest"
ghcr_private: false
```

Note: Replace <YOUR_USER> with your GitHub username.

Step 3: Create deploy-app.yml

infrastructure/ansible/playbooks/deploy-app.yml

```
--  
- name: Deploy Flask App using Docker  
  hosts: all  
  become: yes  
  
vars:  
  image: "{{ image_ref }}"  
  
tasks:  
  - name: Check if docker exists  
    command: bash -lc "command -v docker || true"  
    register: docker_cmd  
    changed_when: false  
  
  - name: Install Docker via script if missing  
    shell: |  
      set -e  
      chmod +x {{ playbook_dir }}/../get-docker.sh || true  
      sudo sh {{ playbook_dir }}/../get-docker.sh  
    when: docker_cmd.stdout == ""  
    args:  
      executable: /bin/bash  
  
  - name: Ensure docker service is started  
    service:  
      name: docker  
      state: started  
      enabled: yes  
  
  - name: Verify Docker service  
    shell: systemctl is-active docker
```

Course Title Deliverable Type

```
register: docker_status
changed_when: false
failed_when: docker_status.stdout.strip() != "active"
```

```
- name: Log in to GHCR (only if private)
  command: >
    docker login ghcr.io
    -u {{ lookup('env','GHCR_USER') }}
    --password {{ lookup('env','GHCR_PAT') }}
  register: ghcr_login
  changed_when: '"Login Succeeded" in ghcr_login.stdout'
  retries: 3
  delay: 5
  until: ghcr_login.rc == 0
  when: ghcr_private | bool
```

```
- name: Pull Flask application image
  command: docker pull {{ image }}
  register: pull_result
  retries: 3
  delay: 5
  until: pull_result.rc == 0
  ignore_errors: no
```

```
- name: Stop existing container if running
  command: docker stop capstone_flask
  ignore_errors: yes
```

```
- name: Remove old container if exists
  command: docker rm capstone_flask
  ignore_errors: yes
```

```
- name: Run application container
```

```
command: >
  docker run -d --name capstone_flask
  -p {{ app_port }}:5000
  {{ image }}

  - name: Wait for Flask application to start
    uri:
      url: "http://localhost:{{ app_port }}/api/health"
      method: GET
      status_code: 200
    register: app_status
    retries: 15
    delay: 4
    until: app_status.status == 200

  - name: Display success message
    debug:
      msg: "☑️ Flask application deployed successfully and is healthy!"
```

This playbook installs Docker (only if missing), optionally logs in to GHCR if your package is private, pulls the image, and runs it.

Task 6: GitHub Actions: CI (build, test, push) & Deploy

In this final task, you will integrate all components, run comprehensive tests, and validate the complete DevOps pipeline.

Step 1: Builds, tests, and pushes the Docker image to GHCR using GITHUB_TOKEN

.github/workflows/ci.yml

```
name: DevOps Capstone - Continuous Integration
```

```
on:
```

```
  push:
```

```
    branches: [ main ]
```

```
  pull_request:
```

```
    branches: [ main ]
```

```
  workflow_dispatch: {}
```

```
env:
```

```
  IMAGE: ghcr.io/${{ github.repository }}/capstone-flask
```

```
  PYTHON_VERSION: '3.11'
```

```
jobs:
```

```
  build-test:
```

```
    runs-on: ubuntu-latest
```

```
    steps:
```

```
      - uses: actions/checkout@v4
```

```
      - uses: actions/setup-python@v4
```

```
        with:
```

```
          python-version: ${{ env.PYTHON_VERSION }}
```

```
      - name: Install deps & run tests
```

```
        working-directory: application/backend
```

```
run: |
  python -m pip install --upgrade pip
  pip install -r requirements.txt
  pytest -v --maxfail=1 --disable-warnings \
    --html=report.html --self-contained-html \
    --cov=app --cov-report=xml --cov-report=html
```

```
- name: Upload test artifacts
  uses: actions/upload-artifact@v4
  with:
    name: test-report
    path: application/backend/report.html
```

```
docker-build-push:
  runs-on: ubuntu-latest
  needs: build-test
  permissions:
    contents: read
    packages: write
  steps:
    - uses: actions/checkout@v4
```

```
- name: Log in to GHCR
  uses: docker/login-action@v3
  with:
    registry: ghcr.io
    username: ${{ github.actor }}
    password: ${{ secrets.GITHUB_TOKEN }}
```

```
- name: Set up Buildx
  uses: docker/setup-buildx-action@v3
```

```
- name: Build & push
```

Course Title Deliverable Type

```
uses: docker/build-push-action@v5
with:
  context: application/backend
  push: true
  tags: |
    ${{ env.IMAGE }}:latest
    ${{ env.IMAGE }}:${{ github.sha }}
```

Note: This uses GITHUB_TOKEN to push to GHCR for this repository namespace. If your package is private and later pulled by Ansible on the same runner, add a login step in deploy workflow or playbook (we added in playbook already, controlled by ghcr_private).

Step 2: Runs Ansible to deploy (either automatically after CI or manually)

.github/workflows/deploy.yml

```
name: DevOps Capstone - Deployment

on:
  workflow_run:
    workflows: ["DevOps Capstone - Continuous Integration"]
    types: [completed]
    branches: [ main ]
  workflow_dispatch:
    inputs:
      ghcr_private:
        description: "Is GHCR package private?"
        required: true
        type: choice
        options: [ "false", "true" ]
        default: "false"
```

jobs:

deploy:

```

runs-on: ubuntu-latest

permissions:
  contents: read
  packages: read

env:
  GHCR_USER: ${{ secrets.GHCR_USER }}
  GHCR_PAT: ${{ secrets.GHCR_PAT }}

steps:
  - name: Checkout
    uses: actions/checkout@v4

  - name: Install Ansible
    run: |
      python3 -m pip install --upgrade pip
      pip3 install ansible

  - name: Render inventory with chosen visibility
    run: |
      sed -i "s/ghcr_private: .*/ghcr_private: ${github.event.inputs.ghcr_private || 'false'}/" infrastructure/ansible/group_vars/all.yml || true

  - name: Run deployment (local deploy on runner)
    working-directory: infrastructure/ansible
    run: |
      ansible-playbook -i inventories/development.yml playbooks/deploy-app.yml

  - name: Smoke test
    run: |
      curl -fsS http://localhost:5000/api/health

```

What this does:

- If CI succeeded on main, this deploys automatically.
- You can also Run workflow manually and choose whether GHCR is private.
- It deploys to the GitHub runner (ephemeral) for demo.

Task 7: Full Run

In this final task, you will integrate all components, run comprehensive tests, and validate the complete DevOps pipeline.

Step 1: Commit & push:

```
cd ~/Desktop/Capstone-CICD
git add .
git commit -m "CI, Docker, Ansible deploy wired up"
git push
```

Step 2: Watch Actions tab:

- CI builds/tests & pushes image
- Deploy runs Ansible, pulls image, starts container, checks /api/health

← DevOps Capstone - Continuous Integration

Fix: merged backend repo into main Capstone repo #1

Re-run all jobs ...

Summary	
Triggered via push 25 minutes ago	Status
optimisticwaqar pushed → 26fac94 main	Success
	Total duration
	1m 16s
	Artifacts
build-test	1
docker-build-push	

ci.yml
on: push

```

graph LR
    A[build-test] --> B[docker-build-push]
    
```

Run details

Usage

Workflow file

Artifacts		
Produced during runtime		
Name	Size	Digest
test-report	3.81 KB	sha256:8a3d1b0ab325e3a28b676518a0e738d98a08e099179fcba...

Capstone Lab Summary: End-to-End CI/CD Pipeline

The objective of this lab was to implement a **complete CI/CD pipeline** that automatically builds, tests, packages, and deploys a Flask web application using modern DevOps tools — **GitHub Actions, Docker, and Ansible** — while following industry best practices for automation and reliability.

Theoretical Flow and Key Concepts

1. Continuous Integration (CI)

Continuous Integration automates the process of building and testing code whenever developers push updates to the main repository.

In this lab:

- GitHub Actions automatically triggered the CI workflow (ci.yml) when code was pushed or when a pull request was created on the main branch.
- The workflow:
- Checked out the repository.
- Installed dependencies in a Python environment.
- Ran Pytest to execute automated unit tests (test_app.py).
- Generated test coverage reports (HTML and XML).
- Built and pushed a Docker image to GitHub Container Registry (GHCR) tagged as latest and by commit.

This ensured every commit was tested and built consistently, guaranteeing the application is always in a deployable state.

2. Continuous Delivery (CD)

Continuous Delivery focuses on automatically deploying applications once they have passed the integration and testing stages.

In this lab:

The deployment workflow (deploy.yml) ran automatically after the CI pipeline completed successfully.

The workflow:

- Installed Ansible in the GitHub runner.
- Used the Ansible playbook (deploy-app.yml) to deploy the containerized application.

Course Title Deliverable Type

- Pulled the Docker image from GHCR.
- Stopped and removed old containers (if any).
- Started a new Docker container using the latest image.
- Performed a health check using the /api/health endpoint.

This automated flow demonstrates how code moves seamlessly from source control → containerization → deployment without manual intervention.

3. Infrastructure as Code (IaC) with Ansible

Ansible was used to define infrastructure configuration and deployment logic declaratively:

- Inventories defined target hosts (development.yml for local deployment).
- Group variables stored reusable app configuration (all.yml).
- Playbooks described each deployment step in YAML, making it reproducible and easy to maintain.
- The playbook also ensured Docker was installed and running before deploying the app.

By using Ansible, deployment became idempotent (safe to re-run multiple times) and platform-independent.

4. GitHub Container Registry (GHCR)

The Docker image was stored in GHCR, serving as a secure and versioned artifact repository. This ensured that deployments always used a known image version, preventing “works on my machine” issues.

Access was managed via GitHub secrets (GHCR_USER, GHCR_PAT) for authentication, making the workflow both secure and automated.

5. Workflow Orchestration

The pipeline demonstrated two interconnected workflows:

1. CI Pipeline (ci.yml)

- Triggered by a code push or pull request.
- Built and tested the application.
- Pushed the Docker image to GHCR.

2. CD Pipeline (deploy.yml)

- Triggered automatically when the CI pipeline succeeded.
- Used Ansible to deploy the app in a container.

This orchestration implemented the Continuous Deployment concept: deploy automatically after successful testing.

6. Automation, Reproducibility, and Visibility

- Every stage was automated (no manual setup required).
- Each workflow's logs were visible in the GitHub Actions tab.
- If something failed (e.g., Docker login, missing image, or health check), the workflow clearly indicated the reason for failure.
- Using GitHub-hosted runners kept the setup lightweight and reproducible.

Lab review

In the CI/CD pipeline implemented, what triggers automated deployment to production?

- A. Every commit to any branch
- B. Manual approval after successful staging deployment
- C. Automatic deployment on main branch
- D. Only scheduled deployments

Correct Answer: B. Manual approval after successful staging deployment

STOP

You have successfully completed this lab.